

Report: Stock Price Prediction using LSTM, Random Forest & XGBoost

Name: Vaishali Deshmukh

Roll no: 250240325071

1. Objective

To predict stock closing prices using historical data with three models:

I'm testing **three different prediction methods** to see which one gives the most accurate guess of tomorrow's stock-market closing price:

1. **LSTM** – a deep-learning model that looks at the last few weeks of prices in order.
2. **Random Forest** – many decision-tree models averaged together.
3. **XGBoost** – lots of small trees built one after another to correct each other's mistakes.

The goal is to identify which model performs best based on accuracy metrics: RMSE, MAE, and R².

1. MAE (Mean Absolute Error)

What it means:

On average, how many rupees our prediction is wrong by.

Example:

If we predicted ₹105, but the actual price was ₹100, then the error is ₹5.

If the **MAE is 1.5**, it means we are usually off by about ₹1.50 per prediction.

2. RMSE (Root Mean Squared Error)

What it means:

Like MAE, but it **gives more weight to big mistakes**.

Example:

If one day we predict ₹110 but the real price is ₹90, the 20-rupee mistake hurts more in RMSE than in MAE.

3. R² Score (R-squared)

What it means:

It tells how well the model is doing overall.

If it's **close to 1**, the model is doing a great job.

If it's **close to 0**, the model is almost guessing.

2. Dataset Overview

- File Used: stocks.csv
- Columns: Date, Symbol, Open, High, Low, Close, Volume
- Stocks Included: Multiple, handled per unique Symbol

3. Data Preprocessing & Feature Engineering

1. Cleaning Steps:

Sorted data by Symbol and Date

Removed rows with missing values after creating lagged features.

2. Lagged Features:

Moving Averages:

MA5, MA10, and MA20 represent 5-day, 10-day, and 20-day moving averages of the Close price.

Purpose: Capture trend and momentum over short, medium, and longer time frames.

Models Used: All (Random Forest, XGBoost, LSTM, Directional LSTM)

3. Price Range Feature

High_Low = High - Low

Represents the daily volatility or price fluctuation.

Helps the model understand how much the stock moved during the day.

Models Used: All

4. Daily Return

Daily_Return = Percentage change in closing price from the previous day.

Daily_Return = df.groupby('Symbol')['Close'].pct_change()

Captures the direction and strength of short-term changes.

Models Used: All

5. Target Variable Creation

Next_Close = Close price of the next day.

This is the main variable we are trying to predict.

Next_Close = df.groupby('Symbol')['Close'].shift(-1)

Models Used: All regression models (RF, XGBoost, LSTM)


These features were used for Random Forest and XGBoost models. For LSTM, the multivariate input sequence included Open, High, Low, Close, and Volume.

4. Model Architectures & Training

Random Forest:

A machine learning model that builds many decision trees and takes the average of their outputs.


- **No loss function or optimizer needed**, as it's not a neural network.
- Trained using features like Open, High, Low, Close, Volume, and moving averages.



XGBoost:

An advanced tree-based model that builds trees one after another, fixing mistakes step by step.

- **Loss Function:** Mean Squared Error (by default).
- Optimized internally using gradient boosting techniques.




LSTM (Long Short-Term Memory):

A deep learning model that learns from sequences (past 20 days of stock prices). It is good at capturing time-based patterns.

- **Loss Function:** Mean Squared Error (MSE).

- **Optimizer:** Adam (a popular and fast optimizer for deep learning).



5. Evaluation Metrics

Used the following

RMSE: Root Mean Squared Error

MAE: Mean Absolute Error

R² Score: Proportion of variance explained by the model

6. Results Summary

Model Name	MAE	RMSE	R-Squared
Random Forest	9.023	42.856	0.9170
LSTM	9.678	41.989	0.9204
XGBoost	14.914	57.304	0.8517

7. Insights & Conclusions

- In this case study, I built and compared three powerful models — **Random Forest**, **XGBoost**, and **LSTM** — to predict the next day's closing stock price using historical data.

- Among all models, **LSTM performed the best**. It had the **lowest big error (RMSE)** and explained the **highest portion of stock price movement (R² = 0.920)**.
- **Random Forest** also performed well, giving similar average accuracy with only slightly higher errors.
- **XGBoost** lagged behind both in terms of accuracy and consistency.

7. Challenges Faced

Large Dataset Processing

The dataset had over **3 lakh records**, which increased training time.

Overfitting Risk

Tree-based models like Random Forest gave **very high R²**, which looked good but were **too close to the naïve baseline**.

Time Formatting Errors

Faced issues with inconsistent date formats (e.g., 13-01-1998) while parsing.

8. Future Enhancements

Adopt Transformer-Based Models

Explore advanced deep learning architectures like the **TimeSeriesTransformer**, which can better capture long-term dependencies and patterns in sequential data.

Enhance prediction accuracy by incorporating **news sentiment analysis, macroeconomic indicators, or sector-specific trends**, which can provide deeper insights into stock price movements.

Conclusion:

In this project, I built models to predict the next day's stock closing price using past data. I tried **Random Forest**, **XGBoost**, and **LSTM** models.

- **LSTM** gave the best results because it understands time patterns well.
- **Random Forest** was fast and easy to use, giving decent accuracy.
- **XGBoost** worked well but didn't beat LSTM.